

Assignment 6

Group 5

Student Name: Zhihao Ren, Chu Gong, Juhi Singhal, Zefu Chen

1.

In [1]:

```
from libsettings import *
import yfinance as yf
from yahoofinancials import YahooFinancials
set_stuff()
```

S&P 500 Part

In [2]:

```
spx_df = yf.download('^GSPC', start='1970-01-01', end='2021-05-01', progress=False)
spx_df.head(10)
```

Out[2]:

	Open	High	Low	Close	Adj Close	Volume
Date						
1970-01-02	0.000	93.540	91.790	93.000	93.000	8050000
1970-01-05	0.000	94.250	92.530	93.460	93.460	11490000
1970-01-06	0.000	93.810	92.130	92.820	92.820	11460000
1970-01-07	0.000	93.380	91.930	92.630	92.630	10010000
1970-01-08	0.000	93.470	91.990	92.680	92.680	10670000
1970-01-09	0.000	93.250	91.820	92.400	92.400	9380000
1970-01-12	0.000	92.670	91.200	91.700	91.700	8900000
1970-01-13	0.000	92.610	90.990	91.920	91.920	9870000
1970-01-14	0.000	92.400	90.880	91.650	91.650	10380000
1970-01-15	0.000	92.350	90.730	91.680	91.680	11120000

In [3]:

```
spx_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 12948 entries, 1970-01-02 to 2021-04-30
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   Open        12948 non-null  float64
 1   High        12948 non-null  float64
 2   Low         12948 non-null  float64
 3   Close       12948 non-null  float64
 4   Adj Close   12948 non-null  float64
 5   Volume      12948 non-null  int64   
dtypes: float64(5), int64(1)
memory usage: 708.1 KB
```

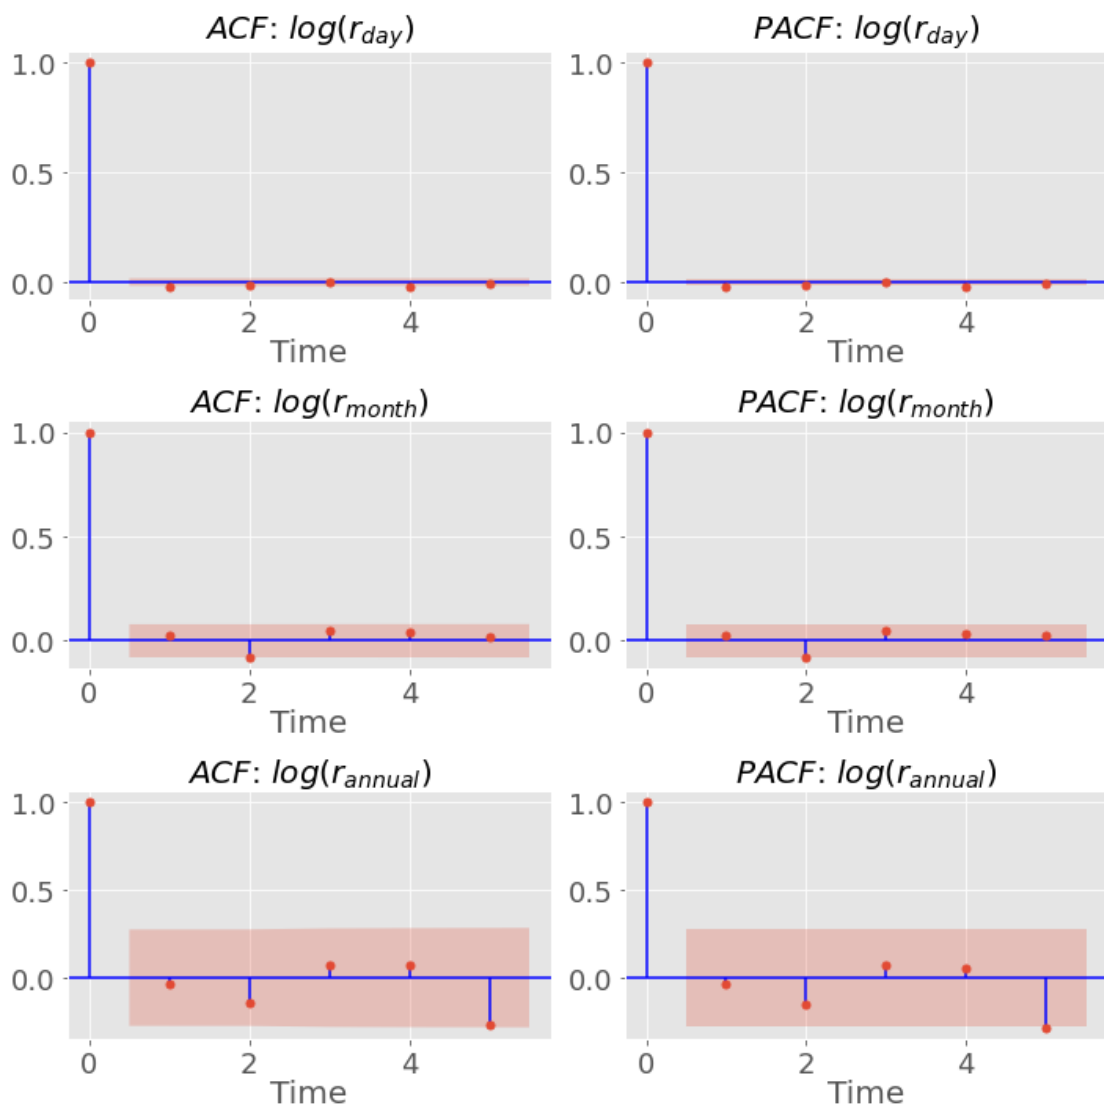
(a)

In [4]:

```
# compute log-return with different freq
ret_daily = np.log(spx_df['Close']).diff().dropna()
ret_monthly = np.log(spx_df.resample('1m').first()['Close']).diff().dropna()
ret_annual = np.log(spx_df.resample('1y').first()['Close']).diff().dropna()

# plot autocorrelation
nlags = 5
fig, axes = plt.subplots(3, 2, figsize=(10, 10))
plot_acf(ret_daily, ax=axes[0][0], lags=nlags)
plot_pacf(ret_daily, ax=axes[0][1], lags=nlags)
plot_acf(ret_monthly, ax=axes[1][0], lags=nlags)
plot_pacf(ret_monthly, ax=axes[1][1], lags=nlags)
plot_acf(ret_annual, ax=axes[2][0], lags=nlags)
plot_pacf(ret_annual, ax=axes[2][1], lags=nlags)

axes[0][0].set_xlabel('Time', fontsize=20)
axes[0][0].set_title("${\\it ACF}$: $log(r_{\\text{day}})$", fontsize=20)
axes[0][1].set_xlabel('Time', fontsize=20)
axes[0][1].set_title("${\\it PACF}$: $log(r_{\\text{day}})$", fontsize=20)
axes[1][0].set_xlabel('Time', fontsize=20)
axes[1][0].set_title("${\\it ACF}$: $log(r_{\\text{month}})$", fontsize=20)
axes[1][1].set_xlabel('Time', fontsize=20)
axes[1][1].set_title("${\\it PACF}$: $log(r_{\\text{month}})$", fontsize=20)
axes[2][0].set_xlabel('Time', fontsize=20)
axes[2][0].set_title("${\\it ACF}$: $log(r_{\\text{annual}})$", fontsize=20)
axes[2][1].set_xlabel('Time', fontsize=20)
axes[2][1].set_title("${\\it PACF}$: $log(r_{\\text{annual}})$", fontsize=20)
fig.tight_layout()
```



(b)

In [5]:

```
# compute average volatility
vol_ann_day = np.std(ret_daily) * np.sqrt(252)
vol_ann_mon = np.std(ret_monthly) * np.sqrt(12)
vol_ann_ann = np.std(ret_annual)

print(f"Average annualized volatility of daily return: {vol_ann_day}")
print(f"Average annualized volatility of monthly return: {vol_ann_mon}")
print(f"Average annualized volatility of annual return: {vol_ann_ann}")
```

```
Average annualized volatility of daily return: 0.1725434114608058
Average annualized volatility of monthly return: 0.15835887231104673
Average annualized volatility of annual return: 0.1586393190851782
```

From above results, it can be showed that the average annualized volatility of daily return is different from the ones computed using monthly return and annual return, which results from the fact that the annualized volatility is an approximation based on the assumption that the data is i.i.d., while i.i.d. assumption cannot be hold in this case.

(c)


In [6]:

```
vol_mon_mon = np.abs(ret_monthly)
vol_mon_day = ret_daily.resample('1m').std() * np.sqrt(21)

fig, ax = plt.subplots()
ax.plot(vol_mon_mon, label='vol of monthly return')
ax.plot(vol_mon_day, label='vol of monthly return based on daily return')
plt.legend()
fig.tight_layout()

print("Mean of vol of monthly return based on daily return: {}".format(vol_mon_day.mean()))
print("Variance of vol of monthly return based on daily return: {}".format(vol_mon_day.var()))
print("Mean of vol of monthly return: {}".format(vol_mon_mon.mean()))
print("Variance of vol of monthly return: {}".format(vol_mon_mon.var()))
print("Correlation of two series: {}".format(vol_mon_day.corr(vol_mon_mon)))
```

Mean of vol of monthly return based on daily return: 0.042535938839692984
Variance of vol of monthly return based on daily return: 0.0006846084042969371
Mean of vol of monthly return: 0.034432962337145184
Variance of vol of monthly return: 0.0009432054133356655
Correlation of two series: 0.41791439108721495



(d)

In [7]:

```
ret_daily_sq = np.power(ret_daily, 2)
realized_vol = np.sqrt(ret_daily_sq.resample('1m').sum())

y = realized_vol.values[1:]
x = realized_vol.values[:-1]
reg = sm.OLS(y, sm.add_constant(x)).fit(cov_type='HC0')
print(reg.summary())
```

OLS Regression Results

```
=====
==
Dep. Variable:          y    R-squared:          0.3
46
Model:                OLS    Adj. R-squared:       0.3
44
Method:              Least Squares    F-statistic:      50.
77
Date:                Mon, 03 May 2021    Prob (F-statistic):  2.93e-
12
Time:                16:44:41    Log-Likelihood:    150
1.7
No. Observations:    615    AIC:                -299
9.
Df Residuals:        613    BIC:                -299
1.
Df Model:             1
Covariance Type:     HC0
=====
==
               coef    std err          z      P>|z|      [0.025    0.97
5]
-----
--
const         0.0175     0.003     5.461     0.000     0.011     0.0
24
x1            0.5878     0.082     7.125     0.000     0.426     0.7
49
=====
==
Omnibus:          670.975    Durbin-Watson:      2.1
67
Prob(Omnibus):    0.000    Jarque-Bera (JB):    60055.4
48
Skew:             4.925    Prob(JB):            0.
00
Kurtosis:         50.398    Cond. No.            3
8.5
=====
==
```

Notes:

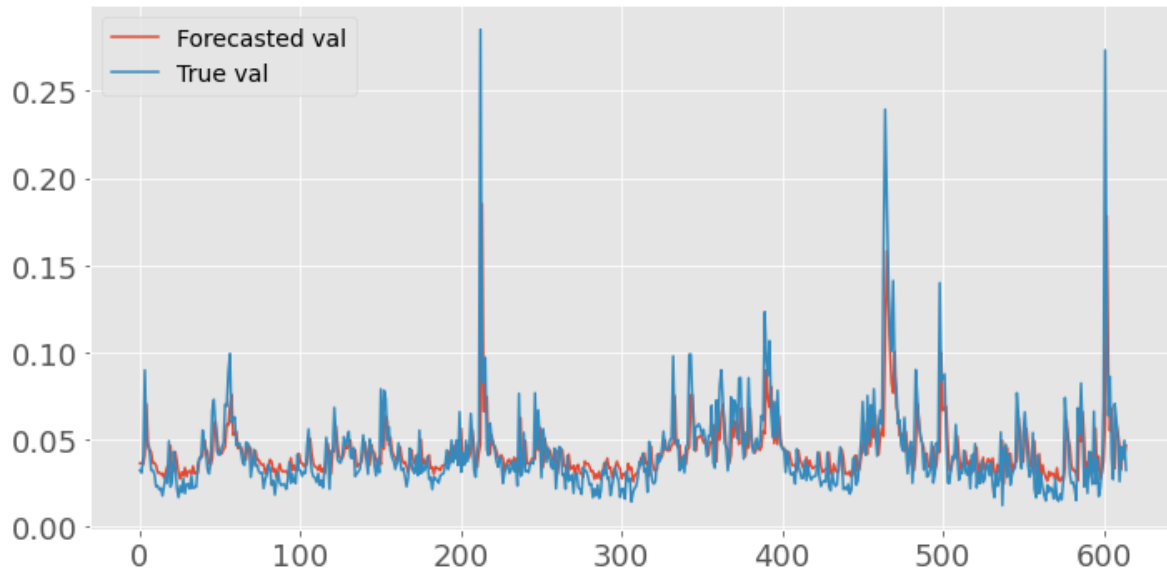
[1] Standard Errors are heteroscedasticity robust (HC0)

From the regression, we can see that the coef is significantly positive (i.e. ~ 0.58), and the R^2 is 0.346. One thing should be noticed is that the kurtosis is 50.398, which suggests a fat tail there.

(e)

In [8]:

```
fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(reg.predict(), label='Forecasted val')
ax.plot(y, label='True val')
plt.legend()
fig.tight_layout()
```



(f)

In [9]:

```
res = np.mean(np.power(reg.resid, 2))
print(f"Mean Squared Error: {res}")
```

Mean Squared Error: 0.00044315856541676763

From the forecast result, it can be seen that the prediction is not good, the norm of the spike cannot be perfectly caught, also the realized vol is more oscillated than the predicted one.

Now we go the same route for GE.

GE Part

In [10]:

```
ge_df = yf.download('GE', start='1970-01-01', end='2021-05-01', progress=False)
ge_df.head(10)
```

Out[10]:

	Open	High	Low	Close	Adj Close	Volume
Date						
1970-01-02	0.776	0.777	0.766	0.767	0.159	2316288
1970-01-05	0.767	0.771	0.757	0.764	0.159	4233216
1970-01-06	0.762	0.762	0.737	0.741	0.154	3544320
1970-01-07	0.744	0.755	0.744	0.745	0.155	4602624
1970-01-08	0.747	0.759	0.747	0.751	0.156	13897728
1970-01-09	0.751	0.752	0.732	0.732	0.152	5940480
1970-01-12	0.732	0.735	0.725	0.731	0.152	4043520
1970-01-13	0.731	0.737	0.730	0.734	0.152	2875392
1970-01-14	0.735	0.752	0.735	0.747	0.155	3694080
1970-01-15	0.747	0.754	0.746	0.749	0.155	2675712

In [11]:

```
ge_df.info()
```

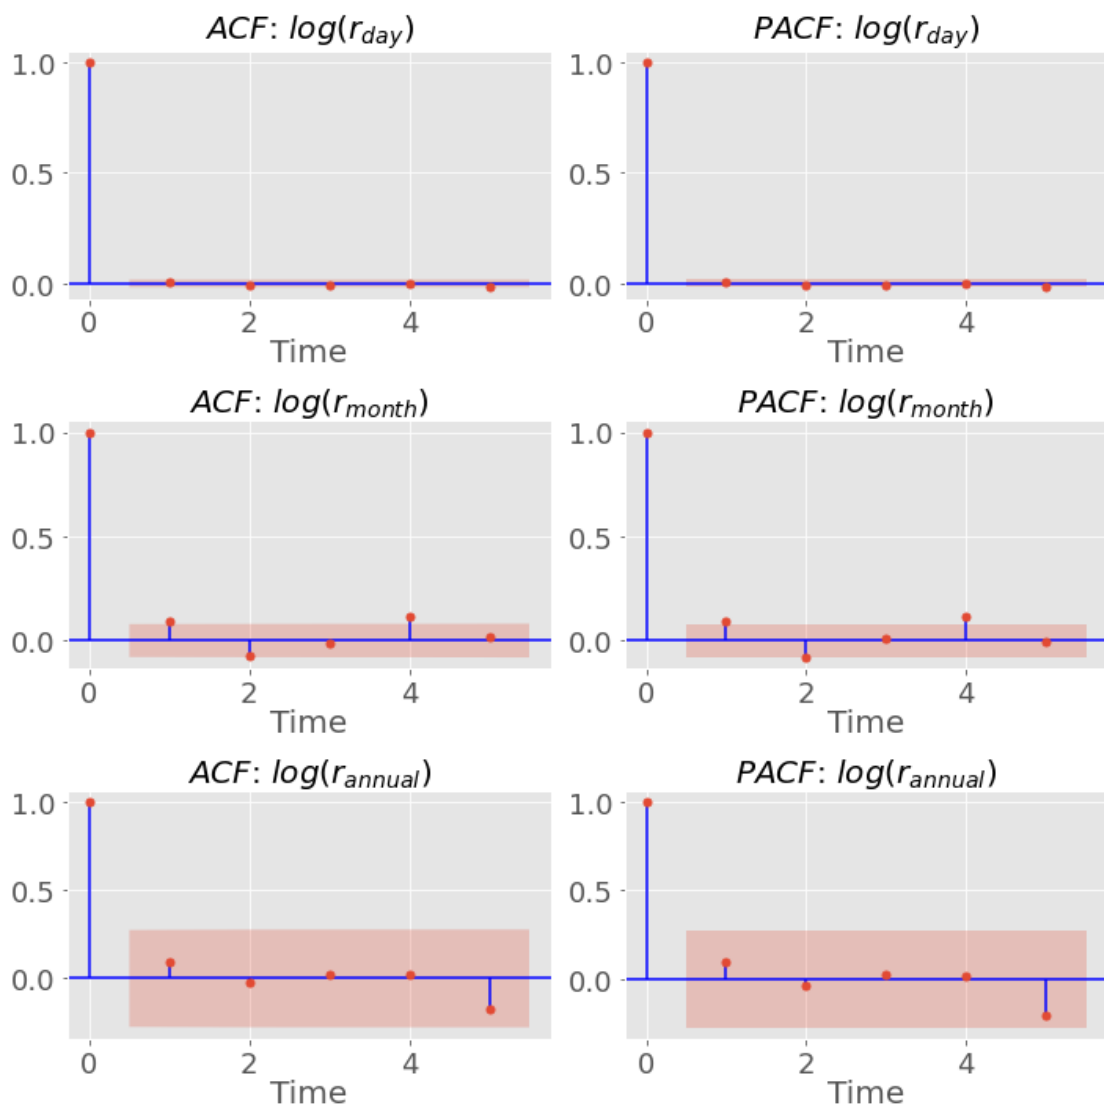
```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 12948 entries, 1970-01-02 to 2021-04-30
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   Open        12948 non-null  float64
 1   High        12948 non-null  float64
 2   Low         12948 non-null  float64
 3   Close       12948 non-null  float64
 4   Adj Close   12948 non-null  float64
 5   Volume      12948 non-null  int64  
dtypes: float64(5), int64(1)
memory usage: 708.1 KB
```


In [12]:

```
# compute log-return with different freq
ret_daily = np.log(ge_df['Close']).diff().dropna()
ret_monthly = np.log(ge_df.resample('1m').first()['Close']).diff().dropna()
ret_annual = np.log(ge_df.resample('1y').first()['Close']).diff().dropna()

# plot autocorrelation
nlags = 5
fig, axes = plt.subplots(3, 2, figsize=(10, 10))
plot_acf(ret_daily, ax=axes[0][0], lags=nlags)
plot_pacf(ret_daily, ax=axes[0][1], lags=nlags)
plot_acf(ret_monthly, ax=axes[1][0], lags=nlags)
plot_pacf(ret_monthly, ax=axes[1][1], lags=nlags)
plot_acf(ret_annual, ax=axes[2][0], lags=nlags)
plot_pacf(ret_annual, ax=axes[2][1], lags=nlags)

axes[0][0].set_xlabel('Time', fontsize=20)
axes[0][0].set_title("${\\it ACF}$: $log(r_{\\text{day}})$", fontsize=20)
axes[0][1].set_xlabel('Time', fontsize=20)
axes[0][1].set_title("${\\it PACF}$: $log(r_{\\text{day}})$", fontsize=20)
axes[1][0].set_xlabel('Time', fontsize=20)
axes[1][0].set_title("${\\it ACF}$: $log(r_{\\text{month}})$", fontsize=20)
axes[1][1].set_xlabel('Time', fontsize=20)
axes[1][1].set_title("${\\it PACF}$: $log(r_{\\text{month}})$", fontsize=20)
axes[2][0].set_xlabel('Time', fontsize=20)
axes[2][0].set_title("${\\it ACF}$: $log(r_{\\text{annual}})$", fontsize=20)
axes[2][1].set_xlabel('Time', fontsize=20)
axes[2][1].set_title("${\\it PACF}$: $log(r_{\\text{annual}})$", fontsize=20)
fig.tight_layout()
```



In [13]:

```
# compute average volatility
vol_ann_day = np.std(ret_daily) * np.sqrt(252)
vol_ann_mon = np.std(ret_monthly) * np.sqrt(12)
vol_ann_ann = np.std(ret_annual)

print(f"Average annualized volatility of daily return: {vol_ann_day}")
print(f"Average annualized volatility of monthly return: {vol_ann_mon}")
print(f"Average annualized volatility of annual return: {vol_ann_ann}")
```

Average annualized volatility of daily return: 0.2831381612138961
Average annualized volatility of monthly return: 0.2691433654414752
Average annualized volatility of annual return: 0.2886513753205717

From above results, it can be showed that the average annualized volatility are slightly different , which results from the fact that the annualized volatility is an approximation based on the assumption that the data is i.i.d., while i.i.d. assumption cannot be hold in this case.

(c)


In [14]:

```
vol_mon_mon = np.abs(ret_monthly)
vol_mon_day = ret_daily.resample('1m').std() * np.sqrt(21)

fig, ax = plt.subplots()
ax.plot(vol_mon_mon, label='vol of monthly return')
ax.plot(vol_mon_day, label='vol of monthly return based on daily return')
plt.legend()
fig.tight_layout()

print("Mean of vol of monthly return based on daily return: {}".format(vol_mon_day.mean()))
print("Variance of vol of monthly return based on daily return: {}".format(vol_mon_day.var()))
print("Mean of vol of monthly return: {}".format(vol_mon_mon.mean()))
print("Variance of vol of monthly return: {}".format(vol_mon_mon.var()))
print("Correlation of two series: {}".format(vol_mon_day.corr(vol_mon_mon)))
```

Mean of vol of monthly return based on daily return: 0.07143836967255703
Variance of vol of monthly return based on daily return: 0.00163472828005728
12
Mean of vol of monthly return: 0.056031956074503536
Variance of vol of monthly return: 0.0029231745718070103
Correlation of two series: 0.45446269609319506



In [15]:

```
ret_daily_sq = np.power(ret_daily, 2)
realized_vol = np.sqrt(ret_daily_sq.resample('1m').sum())

y = realized_vol.values[1:]
x = realized_vol.values[:-1]
reg = sm.OLS(y, sm.add_constant(x)).fit(cov_type='HC0')
print(reg.summary())
```

OLS Regression Results

```
=====
==
Dep. Variable:          y    R-squared:                0.4
87
Model:                  OLS    Adj. R-squared:           0.4
87
Method:                 Least Squares    F-statistic:        18
1.6
Date:                   Mon, 03 May 2021    Prob (F-statistic):    1.91e-
36
Time:                   16:44:43    Log-Likelihood:        130
9.9
No. Observations:       615    AIC:                -261
6.
Df Residuals:           613    BIC:                -260
7.
Df Model:                1
Covariance Type:        HC0
=====
==
               coef    std err          z      P>|z|      [0.025    0.97
5]
-----
--
const          0.0215     0.003     6.592     0.000     0.015     0.0
28
x1             0.6981     0.052    13.477     0.000     0.597     0.8
00
=====
==
Omnibus:            449.691    Durbin-Watson:        2.3
94
Prob(Omnibus):      0.000    Jarque-Bera (JB):      11839.3
40
Skew:               2.906    Prob(JB):              0.
00
Kurtosis:           23.694    Cond. No.              2
5.0
=====
==
```

Notes:


[1] Standard Errors are heteroscedasticity robust (HC0)

From the regression, we can see that the coef is significantly positive (i.e. ~0.698), and the R^2 is 0.487. One thing should be noticed is that the kurtosis is 23.694, which suggests a fat tail there.

(e)

In [16]:

```
fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(reg.predict(), label='Forecasted val')
ax.plot(y, label='True val')
plt.legend()
fig.tight_layout()
```



(f)

In [17]:

```
res = np.mean(np.power(reg.resid, 2))
print(f"Mean Squared Error: {res}")
```

Mean Squared Error: 0.0008271252767787558

From the forecast result, it can be seen that the prediction is not good, the norm of the spike cannot be perfectly caught, also the realized vol is more oscillated than the predicted one.

In []:

PS6-out-of-sample-test

May 5, 2021

0.0.1 Q2

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
import yfinance as yf
import arch
from arch.univariate import ConstantMean, GARCH, Normal
from arch import arch_model
from statsmodels.tsa.arima.model import ARIMA
import statsmodels.api as sm
import warnings
warnings.filterwarnings("ignore")

In [2]: df_sp500_ret = yf.download("^GSPC", start="1970-01-01", end="2021-05-01",
                                progress=False)[["Adj Close"]]
df_ge_ret = yf.download("GE", start="1970-01-01", end="2021-05-01",
                        progress=False)[["Adj Close"]]
```

We will present all question results for S&P500 first, and then GE.

```
In [3]: # monthly log return
df_sp500_ret["yearmonth"] = df_sp500_ret.index.astype(str).str[:7]
df_sp500_ret_monthly = df_sp500_ret.drop_duplicates(subset=["yearmonth"], keep="last")
df_sp500_ret_monthly["log_Close"] = np.log(df_sp500_ret_monthly["Adj Close"])
df_sp500_ret_monthly["log_return"] = 100*(df_sp500_ret_monthly["log_Close"] - \
                                          df_sp500_ret_monthly["log_Close"].shift(1))

In [4]: arr_monthly_ret = df_sp500_ret_monthly["log_return"].dropna()
```

(a) GARCH(1, 1)

```
In [5]: # out-of-sample forecast
split_date = "2014-12-31"

am = ConstantMean(arr_monthly_ret[:split_date])
```

```
# am = ConstantMean(arr_monthly_ret)
am.volatility = GARCH(1, 0, 1)
am.distribution = Normal()
res = am.fit(dis="off")

print(res.summary())
```

```

Constant Mean - GARCH Model Results
=====
Dep. Variable:          log_return      R-squared:          0.000
Mean Model:            Constant Mean    Adj. R-squared:      0.000
Vol Model:             GARCH            Log-Likelihood:     -1551.91
Distribution:          Normal           AIC:                3111.82
Method:                Maximum Likelihood BIC:                3128.99
                                           No. Observations:   540
Date:                  Wed, May 05 2021 Df Residuals:        539
Time:                  21:54:14         Df Model:            1
                               Mean Model
=====
              coef      std err          t      P>|t|    95.0% Conf. Int.
-----
mu            0.6483      0.185        3.506  4.541e-04 [ 0.286,  1.011]
                               Volatility Model
=====
              coef      std err          t      P>|t|    95.0% Conf. Int.
-----
omega         0.8052      0.513        1.568    0.117    [-0.201,  1.812]
alpha[1]       0.1227  3.354e-02        3.657  2.547e-04 [5.693e-02,  0.188]
beta[1]        0.8442  3.468e-02       24.342  7.067e-131 [ 0.776,  0.912]
=====

Covariance estimator: robust

```

From the GARCH(1, 1) estimated results, we noted that $\alpha + \beta$ is close to 1, which means that the volatility of monthly return is highly persistent.


(b) realized monthly volatility and conditional GARCH(1,1) volatility

```
In [6]: arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
arr_realized_vol = np.abs(arr_ret_error)
arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
arr_conditional_vol[split_date] = arr_realized_vol[split_date]

for i in range(1, len(arr_conditional_vol[split_date:])):
    arr_conditional_vol[i] = np.sqrt(res.params["omega"]
    + res.params["alpha[1]"]*arr_realized_vol[i-1]**2
    + res.params["beta[1]"]*arr_conditional_vol[i-1]**2)
```

```
# arr_conditional_vol = res.conditional_volatility
plt.title("out-of-sample")
plt.plot(arr_realized_vol, label="realized volatility")
plt.plot(arr_conditional_vol, label="GARCH(1, 1) conditional volatility")
plt.legend()
```

Out[6]: <matplotlib.legend.Legend at 0x12a7a7cf8>



```
In [7]: model = sm.OLS(arr_realized_vol, sm.add_constant(arr_conditional_vol))
results = model.fit()
results.summary()
```

Out[7]: <class 'statsmodels.iolib.summary.Summary'>
 """

```

                                OLS Regression Results
=====
Dep. Variable:                  log_return    R-squared:                  0.083
Model:                            OLS        Adj. R-squared:              0.071
Method:                 Least Squares    F-statistic:                  6.785
Date:                  Wed, 05 May 2021    Prob (F-statistic):          0.0111
Time:                      21:54:20    Log-Likelihood:              -187.83
No. Observations:                  77        AIC:                        379.7
Df Residuals:                      75        BIC:                        384.3
Df Model:                            1
Covariance Type:                  nonrobust


```


	coef	std err	t	P> t	[0.025	0.975]
const	0.5159	1.031	0.500	0.618	-1.538	2.570
0	0.6469	0.248	2.605	0.011	0.152	1.142
Omnibus:		26.142	Durbin-Watson:			1.552
Prob(Omnibus):		0.000	Jarque-Bera (JB):			38.463
Skew:		1.447	Prob(JB):			4.45e-09
Kurtosis:		4.900	Cond. No.			14.1

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 """

```
In [8]: sm.graphics.plot_fit(results,1)
plt.show()
```



(c) RMSE of GARCH forecasts

```
In [9]: RMSE = np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum() / \
len(arr_conditional_vol))
```

RMSE

Out[9]: 2.9450054362039064

(d) find the best GARCH(p, q) model

```
In [10]: df_model_OS_results = pd.DataFrame(
        index=["BIC", "RMSE", "MAE", "MAPE"],
        columns=["GARCH({}, {})".format(p+1, q+1) for p in range(3)
                for q in range(3)])

In [11]: for p in range(1, 4):
        for q in range(1, 4):
            split_date = "2014-12-31"
            am = ConstantMean(arr_monthly_ret[:split_date])
            # am = ConstantMean(arr_monthly_ret)
            am.volatility = GARCH(p, 0, q)
            am.distribution = Normal()
            res = am.fit(dis="off")

            garch_name = "GARCH({}, {})".format(p, q)
            df_model_OS_results.loc["BIC", garch_name] = res.bic

            # print(res.summary())
            arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
            arr_realized_vol = np.abs(arr_ret_error)
            arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
            # arr_conditional_vol[split_date] = arr_realized_vol[split_date]
            arr_conditional_vol[:max(p, q)] = arr_realized_vol[:max(p, q)]

            for i in range(max(p, q), len(arr_conditional_vol[split_date:])):
                step_sum = res.params["omega"]
                for j in range(p):
                    step_sum += \
                        res.params["alpha[{}]".format(j+1)]*arr_realized_vol[i-1-j]**2
                for k in range(q):
                    step_sum += \
                        res.params["beta[{}]".format(k+1)]*arr_conditional_vol[i-1-k]**2
                arr_conditional_vol[i] = np.sqrt(step_sum)

            df_model_OS_results.loc["RMSE", garch_name] = \
                np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum() / \
                        len(arr_conditional_vol))
            df_model_OS_results.loc["MAE", garch_name] = \
                np.abs(arr_realized_vol - arr_conditional_vol).sum() / \
                len(arr_conditional_vol)
            df_model_OS_results.loc["MAPE", garch_name] = np.abs(
                100*(arr_realized_vol - arr_conditional_vol)/arr_realized_vol).sum() / \
                len(arr_conditional_vol)

In [12]: df_model_OS_results
```

```

Out[12]:      GARCH(1, 1) GARCH(1, 2) GARCH(1, 3) GARCH(2, 1) GARCH(2, 2) GARCH(2, 3) \
BIC          3128.99      3134.4      3140.7      3133.71      3139.24      3145.53
RMSE         2.94501      2.94412      2.94998      3.03346      2.94079      2.94521
MAE          2.46352      2.46693      2.47519      2.51738      2.46135      2.46463
MAPE         259.01       291.81      322.253      284.312      265.055      285.198

      GARCH(3, 1) GARCH(3, 2) GARCH(3, 3)
BIC          3138.17      3143.62      3139.75
RMSE         3.08856      3.0858      3.1196
MAE          2.57498      2.55924      2.52404
MAPE         279.237      276.045      271.708

```

From the BIC perspective, GARCH(1,1) has the smallest BIC. But GARCH(1, 2) has lower forecast error. To keep model parsimonious, we choose GARCH(1,1).

0.0.2 Q3 GJR-GARCH model

(a) GJR-GARCH(1, 1)

```

In [13]: am = arch_model(arr_monthly_ret[:split_date], p=1, o=1, q=1)
res = am.fit(dis="off")
print(res.summary())

```

```

              Constant Mean - GJR-GARCH Model Results
=====
Dep. Variable:          log_return      R-squared:                0.000
Mean Model:             Constant Mean  Adj. R-squared:          0.000
Vol Model:              GJR-GARCH      Log-Likelihood:        -1549.93
Distribution:           Normal          AIC:                    3109.85
Method:                Maximum Likelihood BIC:                    3131.31
                               No. Observations:                540
Date:                  Wed, May 05 2021 Df Residuals:            539
Time:                  21:54:43         Df Model:                1
                               Mean Model
=====
              coef      std err          t      P>|t|  95.0% Conf. Int.
-----
mu           0.5673      0.175       3.244  1.178e-03 [ 0.225, 0.910]
              Volatility Model
=====
              coef      std err          t      P>|t|  95.0% Conf. Int.
-----
omega        1.3454      2.956       0.455    0.649 [-4.449, 7.140]
alpha[1]     0.0536      0.152       0.353    0.724 [-0.244, 0.351]
gamma[1]     0.1045      0.209       0.501    0.616 [-0.304, 0.513]
beta[1]      0.8254      0.135       6.119  9.410e-10 [ 0.561, 1.090]
=====

```

Covariance estimator: robust

(b) realized monthly volatility and conditional GJR-GARCH(1,1) volatility


```
In [14]: arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
arr_realized_vol = np.abs(arr_ret_error)
arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
arr_conditional_vol[split_date] = arr_realized_vol[split_date]

for i in range(1, len(arr_conditional_vol[split_date:])):
    temp_sum = res.params["omega"] \
        + res.params["alpha[1]"]*arr_realized_vol[i-1]**2 \
        + res.params["beta[1]"]*arr_conditional_vol[i-1]**2
    if arr_ret_error[i-1]<0:
        temp_sum += res.params["gamma[1]"]*arr_realized_vol[i-1]**2

    arr_conditional_vol[i] = np.sqrt(temp_sum)

# arr_conditional_vol = res.conditional_volatility
plt.title("out-of-sample")
plt.plot(arr_realized_vol, label="realized volatility")
plt.plot(arr_conditional_vol, label="GJR-GARCH(1, 1) conditional volatility")
plt.legend()

Out[14]: <matplotlib.legend.Legend at 0x11b5c27f0>
```



```
In [15]: model = sm.OLS(arr_realized_vol, sm.add_constant(arr_conditional_vol))
         results = model.fit()
         results.summary()
```

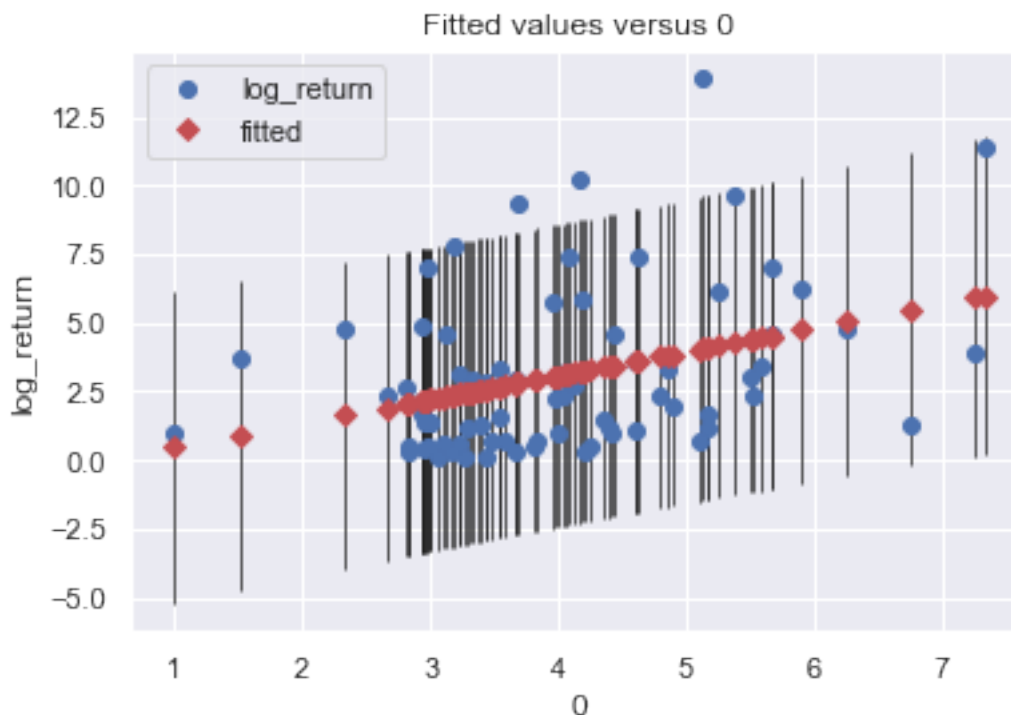
```
Out[15]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:                  log_return    R-squared:                  0.125
Model:                            OLS        Adj. R-squared:              0.113
Method:                 Least Squares    F-statistic:                  10.72
Date:                 Wed, 05 May 2021    Prob (F-statistic):          0.00160
Time:                 21:54:45            Log-Likelihood:              -185.92
No. Observations:                  77        AIC:                        375.8
Df Residuals:                      75        BIC:                        380.5
Df Model:                            1
Covariance Type:                nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const          -0.3960      1.107      -0.358      0.721      -2.601      1.809
0              0.8690      0.265       3.274      0.002       0.340      1.398
=====
Omnibus:                 22.527    Durbin-Watson:              1.632
Prob(Omnibus):            0.000    Jarque-Bera (JB):           30.217
Skew:                     1.321    Prob(JB):                   2.75e-07
Kurtosis:                 4.561    Cond. No.                   15.6
=====
```

Notes:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly spec.
"""
```

```
In [16]: sm.graphics.plot_fit(results,1)
         plt.show()
```



(c) RMSE of GJR-GARCH forecasts

```
In [17]: RMSE = np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum() / \
                        len(arr_conditional_vol))
```

RMSE

```
Out[17]: 2.8625505618637703
```

(d) find the best GARCH(p, q) model

```
In [18]: df_model_OS_results = pd.DataFrame(
        index=["BIC", "RMSE", "MAE", "MAPE"],
        columns=["GJR-GARCH({}, {})".format(p+1, q+1) for p in range(3)
                for q in range(3)])
```

```
In [19]: for p in range(1, 4):
        for q in range(1, 4):
            split_date = "2014-12-31"
            am = arch_model(arr_monthly_ret[:split_date], p=p, o=1, q=q)
            res = am.fit(dis="off")

            garch_name = "GJR-GARCH({}, {})".format(p, q)
            df_model_OS_results.loc["BIC", garch_name] = res.bic
```

```

# print(res.summary())
arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
arr_realized_vol = np.abs(arr_ret_error)
arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
# arr_conditional_vol[split_date] = arr_realized_vol[split_date]
arr_conditional_vol[:max(p, q)] = arr_realized_vol[:max(p, q)]

for i in range(max(p, q), len(arr_conditional_vol[split_date:])):
    step_sum = res.params["omega"]
    for j in range(p):
        step_sum += \
            res.params["alpha[{}]".format(j+1)]*arr_realized_vol[i-1-j]**2
    for k in range(q):
        step_sum += \
            res.params["beta[{}]".format(k+1)]*arr_conditional_vol[i-1-k]**2
    if arr_ret_error[i-1]<0:
        temp_sum += res.params["gamma[1]"]*arr_realized_vol[i-1]**2
    arr_conditional_vol[i] = np.sqrt(step_sum)

df_model_OS_results.loc["RMSE", garch_name] = \
    np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum()/ \
        len(arr_conditional_vol))
df_model_OS_results.loc["MAE", garch_name] = \
    np.abs(arr_realized_vol - arr_conditional_vol).sum()/ \
        len(arr_conditional_vol)
df_model_OS_results.loc["MAPE", garch_name] = np.abs(
    100*(arr_realized_vol - arr_conditional_vol)/arr_realized_vol).sum()/ \
        len(arr_conditional_vol)

```

In [20]: df_model_OS_results

```

Out[20]:      GJR-GARCH(1, 1) GJR-GARCH(1, 2) GJR-GARCH(1, 3) GJR-GARCH(2, 1) \
BIC          3131.31      3132.05      3138.18      3132.05
RMSE          2.80641      2.91919      2.91978      2.90454
MAE           2.21098      2.31297      2.29032      2.28718
MAPE          285.652      257.895      260.705      1219.51

      GJR-GARCH(2, 2) GJR-GARCH(2, 3) GJR-GARCH(3, 1) GJR-GARCH(3, 2) \
BIC          3136.66      3142.33      3133.25      3138.7
RMSE          2.85928      2.85385      3.10443      2.94221
MAE           2.23053      2.2074      2.51539      2.28927
MAPE          659.002      352.983      290.32      357.748

      GJR-GARCH(3, 3)
BIC          3131.34
RMSE          3.05816
MAE           2.39153
MAPE          276.74

```

0.0.3 Q4

```
In [21]: # AR1 models (estimate constant vol)
```

```
y = np.array(np.abs(arr_monthly_ret[:split_date] - arr_monthly_ret[split_date:].mean())
reg = sm.OLS(y[1:], sm.add_constant(y[:-1])).fit()
print(reg.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.029
Model:                  OLS    Adj. R-squared:           0.027
Method:                 Least Squares    F-statistic:        16.17
Date:                   Wed, 05 May 2021    Prob (F-statistic):    6.61e-05
Time:                   21:54:55    Log-Likelihood:        -1344.5
No. Observations:       539    AIC:                  2693.
Df Residuals:           537    BIC:                  2702.
Df Model:                1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                2.7466      0.190     14.494      0.000      2.374      3.119
x1                   0.1705      0.042      4.022      0.000      0.087      0.254
=====
Omnibus:                264.327    Durbin-Watson:           2.027
Prob(Omnibus):           0.000    Jarque-Bera (JB):        1823.572
Skew:                    2.054    Prob(JB):                 0.00
Kurtosis:                11.020    Cond. No.                 6.89
=====
```

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [22]: arr_ret_error = arr_monthly_ret[split_date:] - arr_monthly_ret[split_date:].mean()
arr_realized_vol = np.abs(arr_ret_error)
arr_predicted_vol = pd.Series(reg.predict(sm.add_constant(np.array(arr_realized_vol))),
                              index=arr_realized_vol.index)

plt.title("out-of-sample")
plt.plot(arr_realized_vol, label="realized volatility")
plt.plot(arr_predicted_vol, label="predicted volatility")

plt.legend()
```

```
Out[22]: <matplotlib.legend.Legend at 0x1358307f0>
```

```
In [23]: arr_realized_vol = np.array(arr_realized_vol)
arr_predicted_vol = np.array(arr_predicted_vol)

RMSE = np.sqrt(((arr_realized_vol[1:] - arr_predicted_vol[:-1])**2).sum() / \
                len(arr_realized_vol))

MAE = (np.abs(arr_realized_vol[1:] - arr_predicted_vol[:-1])).sum() / \
      len(arr_realized_vol)

print("RMSE is {}; MAE is {}".format(RMSE, MAE))
```

RMSE is 2.7897498871146706; MAE is 2.154831181506927

Out-of-sample comparison

RMSE: GARCH(1, 1) > GJR-GARCH(1,1) > AR(1)

MAE: GARCH(1, 1) > GJR-GARCH(1,1) > AR(1)

In []:

Same process for the GE return, we can get similar conclusion,

```
In [24]: # monthly log return
df_ge_ret["yearmonth"] = df_ge_ret.index.astype(str).str[:7]
df_ge_ret_monthly = df_ge_ret.drop_duplicates(subset=["yearmonth"], keep="last")
```

```
df_ge_ret_monthly["log_Close"] = np.log(df_ge_ret_monthly["Adj Close"])
df_ge_ret_monthly["log_return"] = 100*(df_ge_ret_monthly["log_Close"] - \
                                         df_ge_ret_monthly["log_Close"].shift(1))
```

```
In [25]: arr_monthly_ret = df_ge_ret_monthly["log_return"].dropna()
```

```
In [26]: # out-of-sample forecast
split_date = "2014-12-31"
am = ConstantMean(arr_monthly_ret[:split_date])
# am = ConstantMean(arr_monthly_ret)
am.volatility = GARCH(1, 0, 1)
am.distribution = Normal()
res = am.fit(dis="off")

print(res.summary())
```

```

                        Constant Mean - GARCH Model Results
=====
Dep. Variable:          log_return      R-squared:                0.000
Mean Model:             Constant Mean  Adj. R-squared:          0.000
Vol Model:              GARCH          Log-Likelihood:         -1775.41
Distribution:           Normal         AIC:                    3558.81
Method:                Maximum Likelihood BIC:                  3575.98
                                           No. Observations:      540
Date:                  Wed, May 05 2021 Df Residuals:            539
Time:                  21:55:16         Df Model:                1
                        Mean Model
=====
                        coef      std err          t      P>|t|  95.0% Conf. Int.
-----
mu                   1.0920      0.270       4.046  5.218e-05 [ 0.563, 1.621]
                        Volatility Model
=====
                        coef      std err          t      P>|t|  95.0% Conf. Int.
-----
omega                2.1702      0.844       2.572  1.010e-02 [ 0.517, 3.824]
alpha[1]             0.1290  3.000e-02       4.300  1.710e-05 [7.019e-02, 0.188]
beta[1]              0.8277  2.946e-02      28.095  1.127e-173 [ 0.770, 0.885]
=====

```

Covariance estimator: robust

```
In [27]: arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
arr_realized_vol = np.abs(arr_ret_error)
arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
arr_conditional_vol[split_date] = arr_realized_vol[split_date]

for i in range(1, len(arr_conditional_vol[split_date:])):
```


```

arr_conditional_vol[i] = np.sqrt(res.params["omega"]
    + res.params["alpha[1]"]*arr_realized_vol[i-1]**2
    + res.params["beta[1]"]*arr_conditional_vol[i-1]**2)

# arr_conditional_vol = res.conditional_volatility
plt.title("out-of-sample")
plt.plot(arr_realized_vol, label="realized volatility")
plt.plot(arr_conditional_vol, label="GARCH(1, 1) conditional volatility")
plt.legend()

```

Out[27]: <matplotlib.legend.Legend at 0x135adfeb8>



```

In [28]: model = sm.OLS(arr_realized_vol, sm.add_constant(arr_conditional_vol))
results = model.fit(dis="off")
results.summary()

```

Out[28]: <class 'statsmodels.iolib.summary.Summary'>
 """

```

                                OLS Regression Results
=====
Dep. Variable:                  log_return    R-squared:                  0.046
Model:                            OLS        Adj. R-squared:             0.033
Method:                         Least Squares   F-statistic:                 3.581
Date:                            Wed, 05 May 2021   Prob (F-statistic):          0.0623
Time:                            21:55:20        Log-Likelihood:              -260.61

```

```

No. Observations:      77    AIC:      525.2
Df Residuals:         75    BIC:      529.9
Df Model:              1
Covariance Type:      nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	3.3479	2.276	1.471	0.145	-1.186	7.881
0	0.4649	0.246	1.892	0.062	-0.024	0.954

```

=====
Omnibus:      33.988    Durbin-Watson:      2.029
Prob(Omnibus): 0.000    Jarque-Bera (JB):      63.106
Skew:         1.685    Prob(JB):      1.98e-14
Kurtosis:     5.884    Cond. No.      25.8
=====

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

In [29]: sm.graphics.plot_fit(results,1)
plt.show()

```



```

In [30]: RMSE = np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum() / \
len(arr_conditional_vol))

```

RMSE

Out[30]: 7.470404489846029

```
In [31]: df_model_OS_results = pd.DataFrame(
        index=["BIC", "RMSE", "MAE", "MAPE"],
        columns=["GARCH({}, {})".format(p+1, q+1) for p in range(3)
                for q in range(3)])

for p in range(1, 4):
    for q in range(1, 4):
        split_date = "2014-12-31"
        am = ConstantMean(arr_monthly_ret[:split_date])
        # am = ConstantMean(arr_monthly_ret)
        am.volatility = GARCH(p, 0, q)
        am.distribution = Normal()
        res = am.fit(dis="off")

        garch_name = "GARCH({}, {})".format(p, q)
        df_model_OS_results.loc["BIC", garch_name] = res.bic

        # print(res.summary())
        arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
        arr_realized_vol = np.abs(arr_ret_error)
        arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
        # arr_conditional_vol[split_date] = arr_realized_vol[split_date]
        arr_conditional_vol[:max(p, q)] = arr_realized_vol[:max(p, q)]

        for i in range(max(p, q), len(arr_conditional_vol[split_date:])):
            step_sum = res.params["omega"]
            for j in range(p):
                step_sum += \
                    res.params["alpha[{}]".format(j+1)]*arr_realized_vol[i-1-j]**2
            for k in range(q):
                step_sum += \
                    res.params["beta[{}]".format(k+1)]*arr_conditional_vol[i-1-k]**2
            arr_conditional_vol[i] = np.sqrt(step_sum)

        df_model_OS_results.loc["RMSE", garch_name] = \
            np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum() / \
                    len(arr_conditional_vol))
        df_model_OS_results.loc["MAE", garch_name] = \
            np.abs(arr_realized_vol - arr_conditional_vol).sum() / \
            len(arr_conditional_vol)
        df_model_OS_results.loc["MAPE", garch_name] = np.abs(
            100*(arr_realized_vol - arr_conditional_vol)/arr_realized_vol).sum() / \
            len(arr_conditional_vol)
```

```
In [32]: df_model_OS_results
```

```
Out[32]:
```

	GARCH(1, 1)	GARCH(1, 2)	GARCH(1, 3)	GARCH(2, 1)	GARCH(2, 2)	GARCH(2, 3)	\
BIC	3575.98	3578.94	3585.23	3581.67	3584.28	3590.57	
RMSE	7.4704	7.47695	7.47421	7.53465	7.47893	7.47641	
MAE	5.5226	5.31234	5.27759	5.58248	5.38772	5.35131	
MAPE	1019.13	917.668	923.507	3690.58	910.527	915.82	

	GARCH(3, 1)	GARCH(3, 2)	GARCH(3, 3)
BIC	3586.62	3587.62	3590.14
RMSE	7.68542	7.62345	7.78692
MAE	5.63636	5.51728	5.62788
MAPE	837.012	895.705	5193.92

```
In [ ]:
```

try GJR GARCH

```
In [33]: am = arch_model(arr_monthly_ret[:split_date], p=1, o=1, q=1)
res = am.fit(dis="off")
print(res.summary())
```

```

Constant Mean - GJR-GARCH Model Results
=====
Dep. Variable:          log_return      R-squared:          0.000
Mean Model:            Constant Mean    Adj. R-squared:      0.000
Vol Model:             GJR-GARCH        Log-Likelihood:     -1773.39
Distribution:           Normal          AIC:                3556.78
Method:                Maximum Likelihood BIC:                3578.24
                                     No. Observations:      540
Date:                  Wed, May 05 2021 Df Residuals:        539
Time:                  21:55:26         Df Model:           1
                                     Mean Model
=====
              coef      std err          t      P>|t|    95.0% Conf. Int.
-----
mu           0.9687      0.268        3.612  3.035e-04 [ 0.443,  1.494]
Volatility Model
=====
              coef      std err          t      P>|t|    95.0% Conf. Int.
-----
omega        3.2419      1.872        1.732  8.336e-02 [ -0.428,  6.911]
alpha[1]     0.0651    4.265e-02        1.526    0.127 [-1.850e-02,  0.149]
gamma[1]     0.0997    7.596e-02        1.313    0.189 [-4.915e-02,  0.249]
beta[1]      0.8147    4.789e-02       17.012  6.655e-65 [ 0.721,  0.909]
=====

```

```
Covariance estimator: robust
```

```

In [34]: arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
arr_realized_vol = np.abs(arr_ret_error)
arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
arr_conditional_vol[split_date] = arr_realized_vol[split_date]


for i in range(1, len(arr_conditional_vol[split_date:])):
    temp_sum = res.params["omega"] \
        + res.params["alpha[1]"]*arr_realized_vol[i-1]**2 \
        + res.params["beta[1]"]*arr_conditional_vol[i-1]**2
    if arr_ret_error[i-1]<0:
        temp_sum += res.params["gamma[1]"]*arr_realized_vol[i-1]**2

    arr_conditional_vol[i] = np.sqrt(temp_sum)

# arr_conditional_vol = res.conditional_volatility
plt.title("out-of-sample")
plt.plot(arr_realized_vol, label="realized volatility")
plt.plot(arr_conditional_vol, label="GJR-GARCH(1, 1) conditional volatility")
plt.legend()

```

Out [34]: <matplotlib.legend.Legend at 0x1356754e0>



```

In [35]: model = sm.OLS(arr_realized_vol, sm.add_constant(arr_conditional_vol))
results = model.fit()
results.summary()

```

```
Out [35]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

OLS Regression Results

```
=====
Dep. Variable:          log_return    R-squared:                0.062
Model:                  OLS          Adj. R-squared:            0.050
Method:                 Least Squares  F-statistic:              4.970
Date:                  Wed, 05 May 2021  Prob (F-statistic):      0.0288
Time:                  21:55:28       Log-Likelihood:           -259.99
No. Observations:      77            AIC:                    524.0
Df Residuals:          75            BIC:                    528.7
Df Model:              1
Covariance Type:       nonrobust
=====
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const          2.4500      2.336      1.049      0.298      -2.203      7.103
0              0.5615      0.252      2.229      0.029       0.060      1.063
=====
```

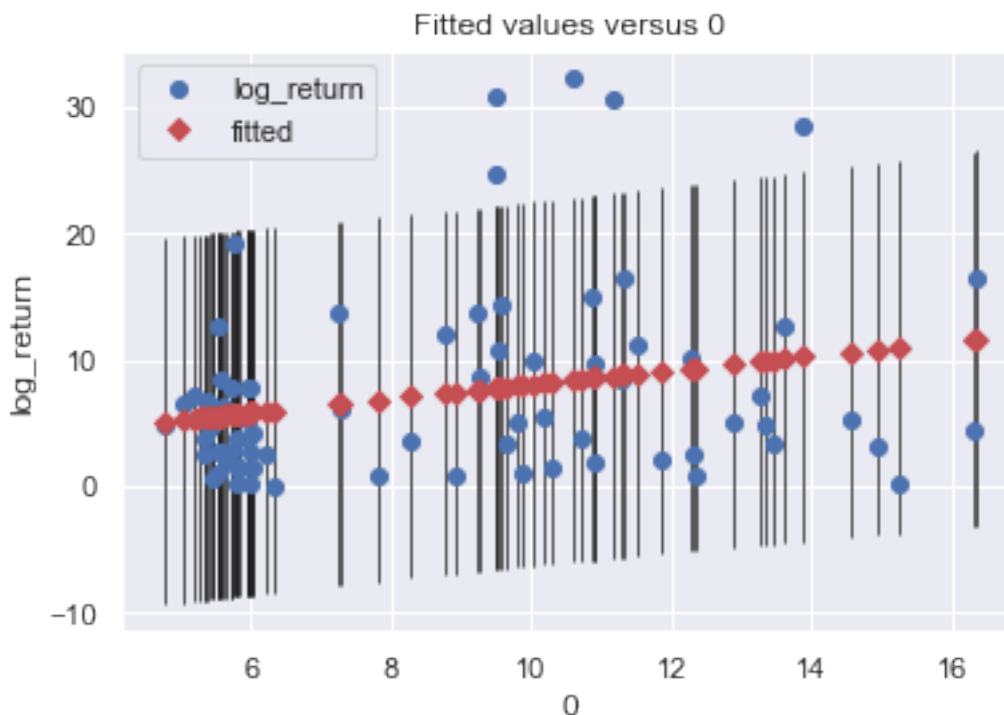
```
=====
Omnibus:                 32.955    Durbin-Watson:              2.082
Prob(Omnibus):            0.000    Jarque-Bera (JB):         60.057
Skew:                    1.640    Prob(JB):                 9.10e-14
Kurtosis:                 5.822    Cond. No.                  26.8
=====
```

Notes:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly spec.
```

```
"""
```

```
In [36]: sm.graphics.plot_fit(results,1)
plt.show()
```

```
In [37]: RMSE = np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum() / \
len(arr_conditional_vol))
```

RMSE

```
Out[37]: 7.350501469111115
```

```
In [38]: df_model_OS_results = pd.DataFrame(
    index=["BIC", "RMSE", "MAE", "MAPE"],
    columns=["GJR-GARCH({}, {})".format(p+1, q+1) for p in range(3)
            for q in range(3)])
```

```
In [39]: for p in range(1, 4):
    for q in range(1, 4):
        split_date = "2014-12-31"
        am = arch_model(arr_monthly_ret[:split_date], p=p, o=1, q=q)
        res = am.fit(dis="off")

        garch_name = "GJR-GARCH({}, {})".format(p, q)
        df_model_OS_results.loc["BIC", garch_name] = res.bic

        # print(res.summary())
        arr_ret_error = arr_monthly_ret[split_date:] - res.params["mu"]
        arr_realized_vol = np.abs(arr_ret_error)
```

```

arr_conditional_vol = pd.Series(index=arr_realized_vol.index)
# arr_conditional_vol[split_date] = arr_realized_vol[split_date]
arr_conditional_vol[:max(p, q)] = arr_realized_vol[:max(p, q)]

for i in range(max(p, q), len(arr_conditional_vol[split_date:])):
    step_sum = res.params["omega"]
    for j in range(p):
        step_sum += \
            res.params["alpha[{}]".format(j+1)]*arr_realized_vol[i-1-j]**2
    for k in range(q):
        step_sum += \
            res.params["beta[{}]".format(k+1)]*arr_conditional_vol[i-1-k]**2
    if arr_ret_error[i-1]<0:
        temp_sum += res.params["gamma[1]"]*arr_realized_vol[i-1]**2
    arr_conditional_vol[i] = np.sqrt(step_sum)

df_model_OS_results.loc["RMSE", garch_name] = \
    np.sqrt(((arr_realized_vol - arr_conditional_vol)**2).sum()/ \
        len(arr_conditional_vol))
df_model_OS_results.loc["MAE", garch_name] = \
    np.abs(arr_realized_vol - arr_conditional_vol).sum()/ \
        len(arr_conditional_vol)
df_model_OS_results.loc["MAPE", garch_name] = np.abs(
    100*(arr_realized_vol - arr_conditional_vol)/arr_realized_vol).sum()/ \
        len(arr_conditional_vol)

```

In [40]: df_model_OS_results

```

Out [40]:      GJR-GARCH(1, 1)  GJR-GARCH(1, 2)  GJR-GARCH(1, 3)  GJR-GARCH(2, 1)  \
BIC          3578.24        3582.84        3589.13        3581.81
RMSE          7.19211        7.25176        7.24878        7.27634
MAE           5.09499        5.08609        5.04614        5.10789
MAPE           530.88        533.241        542.955        384.955

      GJR-GARCH(2, 2)  GJR-GARCH(2, 3)  GJR-GARCH(3, 1)  GJR-GARCH(3, 2)  \
BIC          3588.07        3594.02        3584.86        3585.57
RMSE          7.26603        7.27034        7.66385        7.46572
MAE           5.10192        5.10785        5.30934         5.1266
MAPE           393.536        453.414        459.726        355.085

      GJR-GARCH(3, 3)
BIC          3587.57
RMSE          7.66641
MAE           5.27949
MAPE           482.058

```

In []:

```

In [41]: # AR1 models (estimate constant vol)
y = np.array(np.abs(arr_monthly_ret[:split_date] - arr_monthly_ret[split_date:].mean(

```

```
reg = sm.OLS(y[1:], sm.add_constant(y[:-1])).fit()
print(reg.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          y      R-squared:                0.019
Model:                  OLS    Adj. R-squared:           0.017
Method:                 Least Squares    F-statistic:        10.52
Date:                   Wed, 05 May 2021    Prob (F-statistic):    0.00126
Time:                   21:55:37    Log-Likelihood:       -1597.8
No. Observations:       539    AIC:                 3200.
Df Residuals:           537    BIC:                 3208.
Df Model:                1
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const                4.6207      0.306     15.097      0.000      4.019      5.222
x1                   0.1386      0.043      3.243      0.001      0.055      0.223
=====
Omnibus:                 158.829    Durbin-Watson:           2.030
Prob(Omnibus):            0.000    Jarque-Bera (JB):        410.804
Skew:                     1.475    Prob(JB):                6.24e-90
Kurtosis:                 6.096    Cond. No.                 10.9
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```


In [42]: arr_ret_error = arr_monthly_ret[split_date:] - arr_monthly_ret[split_date:].mean()
         arr_realized_vol = np.abs(arr_ret_error)
         arr_predicted_vol = pd.Series(reg.predict(sm.add_constant(np.array(arr_realized_vol))),
                                       index=arr_realized_vol.index)

         plt.title("out-of-sample")
         plt.plot(arr_realized_vol, label="realized volatility")
         plt.plot(arr_predicted_vol, label="predicted volatility")

         plt.legend()

```

Out[42]: <matplotlib.legend.Legend at 0x135f0b860>



```
In [43]: arr_realized_vol = np.array(arr_realized_vol)
arr_predicted_vol = np.array(arr_predicted_vol)

RMSE = np.sqrt(((arr_realized_vol[1:] - arr_predicted_vol[:-1])**2).sum() / \
               len(arr_realized_vol))

MAE = (np.abs(arr_realized_vol[1:] - arr_predicted_vol[:-1])).sum() / \
      len(arr_realized_vol)

print("RMSE is {}; MAE is {}".format(RMSE, MAE))
```

RMSE is 7.475000329474823; MAE is 4.967676469631952

RMSE: AR(1) > GARCH(1, 1) > GJR-GARCH(1,1)

MAE: GARCH(1, 1) > GJR-GARCH(1,1) > AR(1)

```
In [ ]:
```